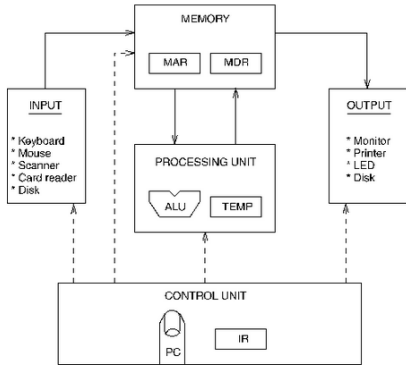


- A **process** (task) is a program in execution, consisting of data, stack, registers needed
- multitasking -- allow several processes coexist (waiting)
 - process table -- stores information about each process each occurrence of a process is called an **instance** think about a timesharing system
 - a computer CANNOT execute 2 processes at the same time

Operating System

- resource allocation
- provides runtime environment
- primary goal is convenience of operation
- secondary goal is efficiency of operation
- a **virtual machine** — it hides complicated and often heterogeneous hardware from the application and provide higher level abstraction to them
- By using CPU scheduling, and virtual memory techniques, an OS can create an illusion of multiple processes, each executing on its own memory and CPU (virtual machine)

The Von Neumann Architecture of Computer Systems



- MAR -- Memory Address Register
- MDR -- Memory Data Register
- ALU -- Arithmetic Logic Unit
- CU -- Control Unit
- CPU -- Central Processing Unit = ALU + CU + Registers + some cache
- I/O -- Input Output Unit
- MU -- Memory Unit
- Von Neumann Computer = CU + ALU + MU + I/O unit

process states

- **new**: the process is being created
- **ready**: process waiting to be assigned to a processor
- **running**: instructions being executed
- **blocked**: waiting for some events to happen (e.g. I/O or reception of signal)
- **terminate**: process has finished execution

Fork()

- used to duplicate a process (2^n)
- The call to fork in the parent return the PID (process ID) of the new child process. The new child process continues to execute just like the original, with the exception that in the child process the call to the fork returns 0

Layered Systems:

- OS → hierarchy layers
 - flexible, has modularity, easier to debug, first layer can be debugged without any concern with the rest of the system
- 5 — operator
4 — user program
3 — i/o management
2 — operator-process communication
1 — memory management
0 — allocation of CPU for processes and multiprogramming

process table (PT) saves Process Control Blocks (PCB)

each process has a PCB with the information:

- process state
- program counter (PC)
- CPU registers
- memory-management information
- accounting information (e.g. time limits, process #, ...)
- I/O status information (e.g. list of open files by the process)

Goals for Scheduling Disciplines

- CPU utilization -- keep CPU as busy as possible
- Fairness -- each process has fair share of CPU
- Minimize response time -- time from the submission of a request until the first response is produced
- Minimize overhead (context swaps)
- Minimize turnaround time -- the interval from the time of submission to the time of completion: waiting time to get into memory + waiting in ready queue + execution on CPU + I/O time
- Maximize throughput -- the number of processes that are completed per unit time

Java

- Objected-oriented, architectural, distributed, multithreaded programming language
- A Java program consists of one or more classes
- Each class ~ architectural byte code
- Runs on Java Virtual Machine (JVM)
- JVM consists of a class loader

Round Robin

if quantum = 20 ms and switching time = 5 ms
% of CPU time 'wasted' = $5 / (20 + 5) = 20\%$

Turnaround time = wait + execution time

Wait time = wait

Threads

- process ~ program in execution + resources needed
- A thread (light weight process (LWP)) is a basic unit of CPU utilization ~ PC + registers + stack space
- a thread shares with peer threads its coded section, data section, and OS resources (such as open files and signals) ~ task
- traditionally heavy weight process = thread + task
- CPU switching much easier
- no memory-management related work need to be done

Monitors

- high-level synchronization,
- consists of procedures, shared resources, and administrative data
- only one process can be active inside a monitor, i.e. only one process can execute any of the methods at any time (mutual exclusion)
- procedures of monitor can only access data inside monitor
- local data of monitor cannot be accessed from outside
- When a task attempts to access a monitor method, it is put in the monitor's entry queue. Each monitor has one waiting queue.
- A **condition variable** is part of the monitor. Sometimes these are called *event queues* or *variable queues*. A condition variable queue can only be accessed with two monitor methods associated with this queue. These methods are typically called **wait** and **signal**. The signal method is called **notify** in Java. Condition variable queues contain processes that have executed a condition variable wait operation. There is one such queue for each condition variable.
- A task that holds the monitor lock may give it up and enter a condition variable queue by executing the corresponding wait method.
- A task that holds the monitor lock may revive a task waiting in a condition variable queue with the notify method of that queue.
- The notify method removes one task from the condition variable queue if the queue is not empty.
- Processes in the monitor queues are waiting to acquire the monitor lock.
- When a task is removed from one of the condition variable queues, it is put in the **waiting queue**.

Weakness:

1. one process active inside monitor => defeats concurrency purpose
 2. nested monitor calls --> deadlock
- Java Synchronization ~ monitor

System Model

- A system consists of a finite number of resources to be distributed among a number of competing processes.
- Resources can only be used by a single process at any single instance of time
- Resources can be of different types, e.g. memory space, CPU cycles, files, I/O devices
- A process must request a resource before using it, and must release the resource after using it
- A normal operation consists of a sequence of events:
 - Request: if request cannot be granted immediately, process must wait until it can acquire the resource
 - Use: operate on the resource (e.g. printing)
 - Release: release the resource

Examples:

- open file, read file, close file
- allocate memory, use memory, free memory
- A set of processes is in a **deadlock** state when every process in the set is waiting for an event that can be caused by another process in the set.

Deadlock modeling

Four necessary conditions for deadlock:

- mutual exclusion -- only one process at a time can use the resource
- hold and wait -- there must exist a process that is holding at least one resource and is waiting to acquire additional resources that are currently being held by other processes.
- no preemption -- resources cannot be preempted; a resource can be released only voluntarily by the process holding it.
- circular wait --
- if graph contains no cycles => no process is in deadlock
- if graph contains cycles => deadlock **may** exist

Deadlock handling strategies

- Ignore the problem altogether
- Prevention -- use a protocol to ensure that the system will never enter a deadlock state
- Detection and recovery -- allow the system to enter a deadlock state and then recover
- Dynamic avoidance -- by careful resource allocation

a) The Ostrich Algorithm

- pretend there's no problem (Windows, UNIX)
- undetected deadlock may result in deterioration of the system performance; eventually, the system will stop functioning and will need to be restarted manually

b) Detection and Recovery

- system monitors requests and releases resources
 - check resource graph to see if cycle exists; kill a process if necessary
 - or if a process has been blocked for, say one hour, kill it
- c) **Deadlock Prevention**
- **mutual exclusion**: printer cannot be shared by 2 processes
 - read-only file -- sharable
- A process never has to wait for a sharable resource
- In general, it is not possible to prevent deadlocks by denying mutual-exclusion condition

• Hold and wait

- have all processes request all their resources before start execution -- inefficient, also, process don't know in advance resources needed before a process can request any additional resources, it must release temporarily all the resources that it currently holds, if successful, it can get the original resources back
 - may have starvation
 - **No preemption**: allows preemption
 - **Circular wait**
- order resource types
process requests resources in an increasing order

d) Deadlock Avoidance

- In prevention strategy, we restrain how request can be made. Here, we want to develop an algorithm to avoid deadlock by making the right choice all the time
- **Banker's Algorithm for Single Resource Type**
- Dijkstra's Banker's Algorithm is an approach to trying to give processes as much as is possible, while guaranteeing no deadlock.
- safe state** -- a state is safe if the system can allocate resources to each process in some order and still avoid a deadlock.
- The banker's algorithm is to consider each request as it occurs, if granting the request would result in a safe state, request is granted otherwise, it is postponed

The benefits of using Threads:

- Performance gains from multiprocessing hardware (parallelism)
- Increased application throughput
- Increased application responsiveness
- Enhanced process-to-process communications
- Efficient use of system resources
- Inherent effectiveness for distributed objects
- Only one binary needed for uniprocessor and multiprocessors
- The ability to create well-structured programs
- There can be a single source for multiple platforms

SJF can be preemptive or nonpreemptive

- nonpreemptive -- when shortest job comes in, it has to wait until the currently-executed job has finished
- preemptive -- when shortest job comes in, currently-executed job will be suspended
- then, shortest-remaining-time-first

Semaphores are not provided by hardware. But they have several attractive properties:

- Machine independent.
- Simple.
- Powerful. Embodiment both exclusion and waiting but do NOT need to be busy waiting
- Correctness is easy to determine.
- Work with many processes.
- Can have many different critical sections with different semaphores.
- Can acquire many resources simultaneously (multiple P's).
- Can permit multiple processes into the critical section at once, if that is desirable.

Semaphores

- Generalization: some code sections may be accessed by a limited number of threads
- A semaphore S is an integer variable, apart from initialization is accessed through standard atomic operations:

Named Semaphores

The advantage of named semaphores is that they provide synchronization between unrelated process and related process as well as between threads. (Related processes refer to parent-child processes.) A named semaphore is created by calling following function:

Unnamed Semaphores

Again, according to the man pages, an unnamed semaphore is placed in a region of memory that is shared between multiple threads (a thread-shared semaphore) or processes (a process-shared semaphore). A thread-shared semaphore is placed in a region where only threads of a process share them, for example a global variable. A process-shared semaphore is placed in a region where different processes can share them, for example something like a shared memory region. An unnamed semaphore provides synchronization between threads and between related processes and are process-based semaphores.

multiple processes

- has states new, ready, running, blocked, terminated
- processes can create child processes
- each process operates independently of the others, has its own PC, stack pointer and address space

multiple threads

- like a process, a thread has states new, ready, running, blocked, terminated
- threads can create child threads
- threads are not independent of each other; threads can read or write over any other's stack

Process Creation

- When a process creates a new process, two possibilities exist in terms of execution
- The parent continues to execute concurrently with its children
- The parent waits until some or all of its children have terminated
- There are two possibilities in terms of the address space of the new process
- the child process is a duplicate of the parent
- the child process has a program loaded into it

Condition Variables

For some problems, using semaphores could be complex. A condition variable is a queue of threads (or processes) waiting for some sort of notifications. Supported by POSIX and SDL; Win-32 events

A condition variable queue can only be accessed with two methods associated with its queue. These methods are typically called wait and signal. The signal method is called notify in Java. Threads waiting for a guard to become true enter the queue. Threads that change the guard from false to true could wake up the waiting threads.

Multiple-process solutions

Lamport's Bakery Algorithm

get a ticket number to purchase baked goods

if two processes pick the number at the same time, it does not guarantee two processes have different #

if number(P_i) < number(P_j) serve P_i first

if number(P_i) = number(P_j), then check i, j (unique process ids); if $i < j$, then serve P_i first

The notation $(a,b) < (c,d)$ is defined as $(a < c)$ or $(a = c \text{ and } b < d)$

Midterm-

RR	run	pri	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	wait	mean	wait
a	6	3																								a	13	9.8
b	4	5																								b	11	
c	1	2																								c	2	
d	3	1																								d	9	
e	7	4																								e	14	

pri	run	pri	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	wait	mean	wait
a	6	3																								a	11	10
b	4	5																								b	0	
c	1	2																								c	17	
d	3	1																								d	18	
e	7	4																								e	4	

FCFS	run	pri	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	wait	mean	wait
a	6	3																								a	0	8.2
b	4	5																								b	6	
c	1	2																								c	10	
d	3	1																								d	11	
e	7	4																								e	14	

SIF	run	pri	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	wait	mean	wait
a	6	3																								a	8	5.4
b	4	5																								b	4	
c	1	2																								c	0	
d	3	1																								d	1	
e	7	4																								e	14	

2) Consider a readers-writers problem with readers-priority, where nr denotes the number of reader threads that have arrived to access the file and nw denotes the

number of writer threads that have arrived. Use guarded commands to:

- present the solution of the readers operations,
- present the solution of the writers operations.

a)

```
void reader()
{
    when ( nw == 0 ) {
        nr++;
    }
    //read
    [nr--];
}
```

b)

```
void writer()
{
    when ( ( nr == 0 ) && ( nw == 0 ) ) {
        nw++;
    }
    //write
    [nw--];
}
```

3) a) In general, a semaphore has two functions, down() which decreases the semaphore value and up(), which increases the semaphore value. What is special about these functions? Give an alternate name that people commonly used for each of the two functions.

b) and c) Two threads, T1() and T2() are accessing the same critical section (CS). They have to wait on two semaphores S1, and S2 before they can enter the CS. Consider the following piece of code for T1 and T2:

T1:

```
down ( S1 );
down ( S2 );
CS();
up ( S2 );
up ( S1 );
```

T2:

```
down ( S2 );
down ( S1 );
CS();
up ( S1 );
up ( S2 );
```

a) A counting semaphore is a semaphore with an integer value that can be larger than 1. A semaphore is an integer variable that, apart from initialization, is accessed only through two standard operations: down() and up(), which differ from usual functions in the way that they are atomic. That is, they must be executed indivisibly. A semaphore is a tool used for synchronization or to achieve mutual exclusion between processes or threads when accessing a critical section.

Alternate names:

```
down() ~ wait(), P()
up() ~ signal, V(), notify()
```

b) Does the code achieve the purpose of mutual exclusion? Why?

Yes, it is because a thread can enter the critical section only if it has acquired the two semaphores ("keys"). If one thread holds a semaphore, the other cannot have it.

c) Is deadlock possible for T1 and T2? Why? If your answer is yes, modify the code so that the two threads are deadlock free. If your answer is no, prove your claim.

Yes. Because threads run concurrently, it could happen that T1 has executed down (S1) and T2 has executed down (S2). So now T1 is waiting for T2 to release (UP) S2 and T2 is waiting for T1 to release S1 and thus they are in deadlock. To prevent deadlock, we can simply require both T1 and T2 to 'lock' (down) semaphores in the same order S1, S2 so that whoever has first locked S1 can proceed to lock S2 and the other thread has to wait on S1. The following is the modified code that is deadlock free:

T1:

```
down ( S1 );
down ( S2 );
CS();
up ( S2 );
```

T2:

```
up ( S1 );
down ( S1 );
down ( S2 );
CS();
up ( S2 );
up ( S1 );
```

- Priority inversion may occur in real-time scheduling — true
- a thread shares with its peers the data section — true
- a process shares with its peers the data section — false
- each processor in symmetric multiprocessing performs self-scheduling — true
- shortest job first scheduling is optimal — true
- A thread can be in only one semaphore's waiting queue at a time — true
- Nonpreemptible resources must be hardware — false
- Processes may deadlock as a result of contending for the processor — false
- An unsafe state is a deadlocked state — false
- Two processes, A and B may deadlock if both of them request three records, 1, 2, and 3 in a database in the order 2, 1, 3 — false

Class Exercises

1. Operating systems manage only hardware. True or false?
 - False

2. What is the primary goal of an OS?

- resource allocation
- provides run-time environment
- primary goal is convenience of operation
- secondary goal is efficiency of operation
- a virtual machine -- it hides complicated and often heterogeneous hardware from the application and provide higher level abstraction to them

3. What limited the size and capabilities of programs in the 1950s?

- Memory

4. What is the difference between a purely layered architecture and a microkernel architecture?

- Not limited to adjacent layer

5. How do microkernels promote portability?

- Different models OS is divided on models

6. What is the output of the following?

```
7. int main()
8. {
9.     int i;
10.
11.     for ( i = 0; i < 3; ++i ) {
12.         fork();
13.         cout << i << endl;
14.     }
15.     return 0;
}
```

Class Exercises

1. What is a semaphore?
 - A semaphore S is an integer variable, apart from initialization is accessed through standard atomic operations
2. A thread can be in only one semaphore's waiting queue at a time. True or false?
 - True. When the semaphore in the queue it goes to sleep and need to be waking up
3. What is a major benefit of implementing semaphores in kernel?
 - It can put it in blocked and do other instead of waiting and sleeping
4. Consider a semaphore that allows the thread with the highest priority to proceed when UP() is called. What potential problem can this cause?
 - If there is someone with low priority and one with high priority. Then the low don't have a chance. "Starvation"
5. What could potentially happen if a thread called UP() without having called the down() operation?
 - It will never get to 0.

Class Exercises

1. Processes do not deadlock as a result of contending for the processor. True or false?
 - True; Processes are preemptive
2. Nonpreemptible resources must be hardware. True or false?
 - False; it is not always hardware

3. Compare and contrast deadlock prevention and deadlock avoidance.

Deadlock Prevention: a set of methods for ensuring that at least one of the necessary conditions for deadlock cannot hold; Prevention by constraining how requests for resources can be made in the system and how they are handled (system design).

Deadlock Avoidance: The system dynamically considers every request and decides whether it is safe to grant it at this point to use during its lifetime. The system requires additional apriori information regarding the overall potential use of each resource for each process. Allows more concurrency.

Similar to the difference between a traffic light and a police officer directing traffic; Preventing deadlocks by constraining how requests for resources can be made in the system and how they are handled (system design).

4. An unsafe state is a deadlocked state. True or false?
 - False

5. Two processes, A and B, each needs three records, 1, 2, and 3 in a database. Is deadlock possible, if both processes request the records in the order 1, 2, 3. How about if they both request the records in the order 2, 1, 3? How about A requests in order 1, 2, 3 and B requests in order 3, 2, 1?

A: 1, 2, 3

B: 1, 2, 3

No possible deadlock because the semaphores in the same order.

A: 2, 1, 3

B: 2, 1, 3

No possible deadlock because the semaphores in the same order.

A: 1, 2, 3

B: 3, 2, 1

Possible deadlock may occur. Threads run concurrently, it could happen that A executed 1 and 2 and B executed 3 then both will be locked waiting for the other thread to release the key. To prevent this both have to be in the same order.

Ex: 1, 2, 3, 1, 4, 2, 1, 5, 6, 2, 1, 3, 7, 6, 3, 2, 1, 3, 3, 7

Page fault: shorts clinic and every time you load a page that's above it's considered a page fault. Only exception is devices ex. 7, 7, 1 frame 1, 12, 13 [2 frame 1]

LRU (Least Recently Used): look back until you know which 1 frame you are replacing

First In, First Out: Short requests @ the top of your stack and keep moving down until you reach the end and start over.

Optimal Replacement: look ahead to see which n-1 frames you are going to keep.

2. (50 points)

a) A computer system has sixteen drives, with 9 processes competing for it need a maximum of two drives. Is the system deadlock free? Give your rea with the help of **resource graph diagrams**.

Solutions:

Yes, the system is deadlock free. This is because in the worst case, all 9 processes request two drives. As we have 16 drives, some of the processes can get two drives to finish their tasks and return the the drives. Then some of the remaining processes can get two drives, finish and return and so on. So eventually, all processes can finish and return all drives.

The figure on the right illustrates the worst case situation. The processes colored in red are competing for a drive but when any green process has finished using and returned the two drives, they can be allocated to the red processes.

a. Contiguous: 401

Because each read and write block has to moved over once, so $200 \times 2 +$ the one extra block.

Linked: 1

Just write the new block making it point to the next block. Update the first block pointer in memory.

Indexed: 1

Just write the new block and update the index in memory.

b. Contiguous: 201

100 blocks (the second half) must be shifted over one block and the new block must be written. As before, the shift takes one read operation and one write operation.

Linked: 102

I must read 100 blocks to find the middle. Then I must write my new block somewhere with the next block pointing to the block after the 100th block. Then I must write the 100th block to point to this new block.

Indexed: 1

Just write the new block and update the index in memory.

c. Contiguous: 1

Just write the new block.

Linked: 3

(assuming that in order to modify the a block's next block pointer, I must first read the whole block in, then write the whole block out just with that pointer changed). First I read the last block as given by the last block pointer, then I write my new block, then I write the last block back modifying its next block pointer to point to my new block, then I update my last block pointer in memory.

Indexed: 1

Just write the new block and update the index in memory.

d. Contiguous: 0

Just point the first block pointer to the second block.

Linked: 1

Read in the first block to get the second block's pointer and then set the first block pointer to point to the second block.

Indexed: 0

Just remove the block from the index in memory.

e. Contiguous: 198

Assuming that we are removing the 101st block, we will have to move 99 blocks one block over this takes a read and a write operation each.

Linked: 102

It takes 101 reads to find the pointer to the 102nd block, then we need to update the 100th block's next block pointer with this value.

Indexed: 0

Just remove the block from the index in memory.

f. Contiguous: 0

Just update the length information stored in memory.

Linked: 199 or 200

We need to update the last block pointer with the second to last block and the only way to get the second to last block is to read the preceding 199 blocks. Then we probably want to mark the 199th block (the new last block) as having a null pointer and this would give the 200th operation. We save the new last block in our last block pointer in memory.

Indexed: 0

Just remove the block from the index in memory.

I/O Systems:

Interrupts

- non-maskable interrupts (NMI) — cannot be masked off, reserved for serious errors
- maskable interrupt — can be turned off by the CPU before the execution of critical instruction sequences that must not be interrupted

Direct Memory Access (DMA)

- Overview: DMA is an operational transfer mode which allows data transfer within memory or between memory and I/O device without processor's intervention. A special DMA controller manages that data transfer.

Disk Scheduling:

- requests are queued
- FCFS (first-come-first-served) scheduling
- SSTF (Shortest-seek-time-first) scheduling
- SCAN scheduling
- the disk arm starts at one end of the disk, and moves toward the other end, servicing requests as it reaches each cylinder; when it reaches the other end, the direction of head movement is reversed,
- e.g. queue = 98, 183, 37, 122, 14, 124, 65, 67
- current head position at 53, moving towards track 0
- service sequence: 37, 14, 65, 67, 98, 122, 124, 183
- C-SCAN (circular scan) scheduling
- when it reaches the other end, it immediately return to the beginning
- service sequence of above example: 65, 67, 98, 122, 124, 183, 14, 37

RAIDS:

- RAID 0: non-redundant block-level striping (improves speed but no fault tolerance)
- RAID 1: mirrored disks (fault tolerance but no improvement in speed)
- RAID 2: memory-style error correcting (ECC), not used in practice because level 3 is better
- RAID 3: bit-level striping with parity (both speed and fault tolerance)
- RAID 4: block-level striping with parity (both speed and fault tolerance)