

Four necessary conditions for deadlock: mutual exclusion, hold and wait, no preemption, circular wait

for every reusable R_i

$t_i = |R_i| = \# \text{ of units of resource of type } i$
 ≥ 0

for every consumable R_i

there exists subset $R \subset P$

producers of R_i

A **bi-partite graph** (bigraph) is a graph whose nodes can be divided into two disjoint sets such that two adjacent nodes cannot come from the same set.

for every reusable resource R_i :

- # of assignment edges $\#(R_i, *) \leq t_i$
- $a_i = t_i - \#(R_i, *)$
- a process cannot request more than the total number of units of a reusable resource
 $\#(P_j, R_i) \leq t_i - \#(R_i, P_j)$

for every consumable resource R_i

- (R_i, P_j) exists iff P_j is a producer of R_i
- $a_i \geq 0$

if $\#(P_i, R_j) > a_j \Rightarrow p_i$ blocked (more requests than available resources)

If a graph can be reduced by all the processes, then there is **no** deadlock.

If a graph cannot be reduced by all the processes, the irreducible processes constitute the set of deadlock processes in the graph

Definition

A **knot** K in a graph is a nonempty set of nodes such that for every node x in K , all nodes in K and only the nodes in K are reachable from x .

(all x , all $y \in K \Rightarrow x \rightarrow y$) AND (all $x \in K$, some z , such that $x \rightarrow z \Rightarrow z \in K$)

A cycle is a necessary condition for a deadlock.

If the graph is expedient, then a knot is a sufficient condition for deadlock.

Systems with Single-unit Requests

A process can only request one resource unit at a time.

Theorem: A knot is a necessary and sufficient condition for deadlock.

Systems with Single-unit Resources

There is only one unit of every resource.

Theorem: A cycle in an *expedient graph* is a necessary and sufficient condition for deadlock.

System with only consumable resources

such as interrupts, signals, messages, and data in I/O buffers

deadlock may occur if a **receive** operation is blocking

Definition: **claim-limited graph** of a consumable resource system:

each resource has 0 available units

it has a request edge (p_i, R_j) iff P_i is a consumer of R_j

- **Basic Client Server Model**
- Characteristics:
 - There are processes offering services (servers)
 - There are processes that use services (clients)
 - Clients and servers can be on different machines
 - Clients follow request/reply model when using services
 - Synchronous communication: request-reply protocol
 - In LANs, often implemented with a connectionless protocol (unreliable)
 - In WANs, communication is typically connection-oriented TCP/IP (reliable)
 - - High likelihood of communication failures

Decentralized Architectures

- In the last couple of years we have been seeing a tremendous growth in peer-to-peer systems.
- Structured P2P: nodes are organized following a specific distributed data structure
- Unstructured P2P: nodes have randomly selected neighbors
- Hybrid P2P: some nodes are appointed special functions in a well-organized fashion

Processor: Provides a set of instructions along with the capability of automatically executing a series of those instructions.

Thread: A minimal software processor in whose context a series of

- instructions can be executed. Saving a thread context implies stopping

- the current execution and saving all the data needed to continue the

execution at a later stage.

Process: A software processor in whose context one or more threads may be

- executed. Executing a thread, means executing a series of instructions

in the context of that thread.

Context Switching:

User-space solution

All operations can be completely handled within a single process =>

implementations can be extremely efficient.

All services provided by the kernel are done **on behalf of the process in**

which a thread resides => if the kernel decides to block a thread, the

entire process will be blocked.

Threads are used when there are lots of external events: **threads block**

on a per-event basis => if the kernel cannot distinguish threads, how can it

support signaling events to them?

Kernel solution

The whole idea is to have the kernel contain the implementation of a thread

package. This means that all operations return as system calls

Operations that block a thread are no longer a problem: the **kernel**

schedules another available thread within the same process.

Handling external events is simple: the **kernel** (which catches all events)

schedules the thread associated with the event.

The big problem is the **loss of efficiency** due to the fact that each thread

operation requires a trap to the kernel.

Virtualization is becoming increasingly important:

Hardware changes faster than software

Ease of portability and code migration

Isolation of failing or attacked components

Basic model

A server is a process that waits for incoming service requests at a

specific transport address. In practice, there is a one-to-one mapping

between a port and a service.

Type of servers:

Iterative servers: handles a request itself; can handle only one client at a time,

Concurrent servers: Does not handle a request itself; pass it to a separate thread or another process

Superservers: Servers that listen to several ports, i.e., provide several

independent services. In practice, when a service request comes

in, they start a subprocess to handle the request (UNIX inetd)

Stateless servers

Never keep accurate information about the status of a client after having

handled a request:

Don't record whether a file has been opened (simply close it again after

access)

Don't promise to invalidate a client's cache

Don't keep track of your clients

Consequences

Clients and servers are **completely independent**

State inconsistencies due to client or server crashes are reduced

Possible loss of performance because, e.g., a server cannot anticipate

client behavior (think of prefetching file blocks)

Software Agents

What is an Agent?

Definition: An autonomous process capable of reacting

to, and initiating changes in its environment, possibly

in collaboration with users and other agents

collaborative agent: collaborate with others in a multiagent system

mobile agent: can move between machines

interface agent: assist users at user-interface level

information agent: manage information from physically different sources

Transport Layer

Important

The transport layer provides the actual communication facilities for

most distributed systems.

Standard Internet Protocols:

- TCP: connection-oriented, reliable, stream-oriented communication
- UDP: unreliable (best-effort) datagram communication

Basic RPC operation

Observations

Application developers are familiar with simple procedure model

Well-engineered procedures operate in isolation (black box)

There is no fundamental reason not to execute procedures on separate machine

Basic principle

Every machine has a timer that generates an interrupt H times per second.

There is a clock in machine p that ticks on each timer interrupt.

Denote the value of that clock by $C_p(t)$, where t is UTC time.

Ideally, we have that for each machine p , $C_p(t) = t$, or, in other words, $dC/dt = 1$.

Real distance is

$$d_i = c\Delta_i - c\Delta_r = \sqrt{(x_i - x_r)^2 + (y_i - y_r)^2 + (z_i - z_r)^2}$$

Lamport's Logical Clock

Condition requirements:

1. for any two events **a** and **b** in a process P_i , if **a** occurs before **b**, then
2. $C_i(a) < C_i(b)$
3. if **a** is the event of sending a message **m** in P_i and **b** is the event of receiving the same message **m** at process P_j , then $C_i(a) < C_j(b)$

Implementation rules:

1. two successive events in P_i
2. $C_i = C_i + d$ ($d > 0$)
3. if **a** and **b** are two successive events in P_i and $a \rightarrow b$ then
4. $C_i(b) = C_i(a) + d$ ($d > 0$)
5. event **a**: sending of message **m** by process P_i , timestamp of message **m**: $t_m = C_i(a)$ then
6. $C_j = \max(C_j, t_m + d)$ $d > 0$

Vector Clocks

Implementation Rules:


1. two successive events **a**, **b** in process P_i :
2. $C_i(b)[i] = C_i(a)[i] + d$ ($d > 0$)
3. event **a** at P_i sending message **m** to process P_j with receiving event **b**; vector timestamp $t_m = C_i(a)$ is assigned to **m**; on receiving **m**, P_j updates C_j as follows:
4. all k , $C_j(b)[k] = \max(C_j(b)[k], t_m[k])$

Assertion.


At any instant

all i , all j : $C_i[i] \geq C_j[i]$

Minimum is 1



Insert 18




Overlay Network

Nodes are formed by the processes of the network.

Structured P2P Systems

Basic idea

- Organize the nodes in a structured overlay network such as a logical ring, and make **specific nodes responsible for services** based only on their ID.
- A common approach is to use a distributed hash table (DHT) to organize the nodes.
- Traditional hash functions convert a key to a hash value, which can be used as an index into a hash table:
 - Keys are unique -- each represents an object to store in the table;
 - The hash function value is used to insert an object in the hash table and to retrieve it.
- In a DHT, data objects and nodes are each assigned a key which hashes to a random number from a very large identifier space (to ensure uniqueness).
- A mapping function assigns objects to nodes, based on the hash function value.
- A lookup, also based on hash function value, returns the network address of the node that stores the requested object.



Note

The system provides an operation **LOOKUP(key)** that will efficiently route the lookup request to the associated node.


Achieved by organizing processes through a distributed hash table (DHT)

e.g. a data item with key k is mapped to the node with the smallest identifier $id \geq k$. This node is referred to as the successor of the datum with key k , and is denoted as $succ(k)$.


Other example

Organize nodes in a d-dimensional space and let every node take the responsibility for data in a specific region. When a node joins => split a region.

An example of reduction of an RAG (Resource Allocation Graph)



Example of RAG with loop but no deadlock



Example of Resource Allocation Graph with loop but no deadlock